



Systems and Technology Group

Using the SDK Sample Programs

Course Code: L2T2H1-11
Cell Ecosystem Solutions Enablement

Course Code: L2T2H1-11 Using the SDK Sample Programs

5/17/2006

© 2005 IBM Corporation

Class Objectives – Things you will learn

- **Examine the contents of the SDK in details, how it was organized, and the software provided in such directories as Samples, Tests, Tools, Workloads, and Application-oriented samples**
- **Analyze the complex multiplication code to understand how DMA and double buffering were used**
- **Learn some tips and techniques to better use the SDK**

Class Agenda

- **The Software Development Kit**
 - SDK contents
 - Samples
 - Tests
 - Tools
 - Workloads
 - Application-oriented code samples
- **A cell programming example**
 - SPE programming key points
 - Complex multiplication – code example
- **SDK tips and techniques**

Trademarks - Cell Broadband Engine™ is a trademark of Sony Computer Entertainment, Inc.

The Software Development Kit

SDK Contents

- **The SDK source code is organized into the following categories:**
- **samples (\$SDK_TOP/src/samples)**
 - simple and concise code examples to demonstrate specific functions, use of tools, libraries, and/or HW features
- **tests (\$SDK_TOP/src/tests)**
 - self-verifying tests use to assure standards compliance, validate libraries and tools
- **tools (\$SDK_TOP/src/tools)**
 - utilities used to generate content or ease programming burden
- **workloads (\$SDK_TOP/src/workloads)**
 - code samples used to characterize the performance of the architecture
- **lib (\$SDK_TOP/src/lib)**
 - libraries and reusable header files

Samples

- **cesof (CBE™ Embedded SPU Object Format)**
 - sample code to demonstrate the object format used to embed SPU objects into PowerPC binaries
- **DMA**
 - sample code to demonstrate non-trivial DMA calls
- **resample**
 - audio resampling code for SP/DP monotonic/stereo audio samples
- **simpleDMA**
- **spu_clean**
 - sample SPU program that clears the register file and local store (including itself)
- **spu_entry**
 - sample crt0 – initializes the stack and stack pointer; calls main; returns main's return value to a controlling PU program in an ABI compliant fashion (exit function).
- **spu_interrupt**
 - sample first level interrupt handler and second level interrupt handler registration function. Demonstration second level decremented interrupt handler.
- **spulet**
 - C-library functions made to run on SPU (printf(), read(), etc.)
- **sync**
 - conditional wait, mutex, and atomic operation sample code
- **tutorial**
 - contains some of the source code used within the tutorial document

Tests

- **abi**
 - set of tests used to validate conformance to the SPU and BE ABI standards
- **asm**
 - set of tests used to verify assembler support of all instructions, parameter forms, and parameter ranges
- **events**
 - set of tests used to validate and demonstrate the handling of user-defined SPU events
- **intrinsic**s
 - set of tests used to validate all VMX and SPU intrinsic
- **lib**
 - suite of self-validating tests used to verify correct operation of the libraries

Tools

- **callthru**
 - callthru source code
- **idl**
 - IDL compiler tool reads a high-level specification describing an interface to a SPU function
 - produces special stub functions to implement the interface in C
 - stubs allow the PU and SPU to communicate through what appear to be ordinary, local procedure calls or method invocations
- **oprofile (in progress)**
 - system-wide profiler for Linux
 - kernel driver and daemon for collecting data
 - several post-profiling tools

Workloads

- **FFT16M**
 - hand-tuned program performing 4-way SIMD SP complex FFT of 16M elements
- **matrix_mul**
 - workload calculates $C = A * B$ where A, B, and C are N x N squared matrices comprised of SP floats.
 - uses block-partitioning algorithm to reduce bandwidth (block size fixed to 64)
- **oscillator**
 - workload used to synthesize two stereo sound files
- **vse_subdiv**
 - workload demonstrating subdivision using contours of variable sharpness
 - displays result in OpenGL output window

Application-oriented Code Samples

- **C**
 - SPE-only library containing functions typically found in standard C99 library
 - includes functions executed by the SPE natively, functions initiated by the SPE but executed by the PPC, and SPE local store functions
 - provides or enhanced common high-level programming functionality
- **audio resample**
 - provides sample audio resampling functions that include
 - monophonic and stereophonic audio
 - unsigned short or FP samples
 - SP and DP computation
- **curves and surfaces**
 - support routines for evaluating quadratic and cubic Bezier curves as well as biquadric/bicubic Bezier surfaces and curved point-normal triangles
- **FFT**
 - highly tuned 1-D FFT as well as base kernel functions that can be used to implement 2-D FFTs

Application-oriented Code Samples (cont.)

- **game math**
 - set of routines implemented with the notion that precision and mathematical accuracy can at times be sacrificed for performance
- **image**
 - includes routines for various size convolutions as well as generation of histograms of byte data
- **large matrix**
 - various utility functions that operate on large vectors/matrices of SP FP numbers
 - size of input vectors and matrices limited by SPE local storage size (no matrix partitioning)
- **math**
 - general purpose math routines tuned to exploit SIMD features
 - most only support SP
 - intended to mimic standard math library functions

Application-oriented Code Samples (cont.)

- **matrix**
 - utility library to operate on matrices and quaternions including inversion, identity, perspective projection, and multiplication
- **misc**
 - routines that do not logically fit into other categories (min, max, rand, clamp, etc.)
- **multi-precision math**
 - performs mathematical functions on unsigned integers of a large number of bits
- **noise**
 - 1-D, 2-D, 3-D, and 4-D noise
 - lattice and non-lattice noise
 - turbulence

Application-oriented Code Samples (cont.)

- **oscillator**
 - two oscillator libraries to create a synthetic environment of configurable directional microphones, a large number of oscillators moving along defined paths, all relative to static microphones
 - computes time delays, volume changes, doppler effects
- **sim**
 - services useful to the full system simulator such as callthru
- **sync**
 - libraries making use of the load-with-reservation and store-conditional functions within CBEA
 - atomic operations, mutexes, conditional variables, and completion variables
 - sample code included in samples dir
- **vector**
 - 15 general purpose routines to operate on vectors

A Cell Programming Example

SPE Programming Key Points

- **SPE memory architecture**
 - each SPE has its own “flat” local memory
 - management of this memory is by explicit software control
 - code and data are moved into and out of this memory by DMA
 - programming the DMA is explicit in SPE code (or in PPC code)
- **Many DMA transactions can be “in flight” simultaneously**
 - e.g. each SPE can have 16 simultaneous outstanding DMA requests
 - a DMA request can be a list of DMA requests
 - DMA latencies can be hidden using multiple buffers and loop blocking in code
 - in contrast to traditional, hierarchical memory architectures that support few simultaneous memory transactions
- **Implications for programming the BE**
 - applications must be partitioned across the processing elements, taking into account the limited local memory available to each SPE

Complex Multiplication – code example

- **(See complex multiplication full code example)**
- **Overall code uses PPC to initiate operation as well as transmit data**
- **PPC waits for SPE to finish**
- **SPE operations are double-buffered**
 - i^{th} DMA operation is started before main loop
 - $(i+1)^{th}$ DMA operation is started at beginning of main loop at secondary storage area
 - main loop then waits on i^{th} DMA to complete
- **Notice that all storage areas are 128B aligned**

Complex Multiplication

- In general, the multiplication of two complex numbers is represented by

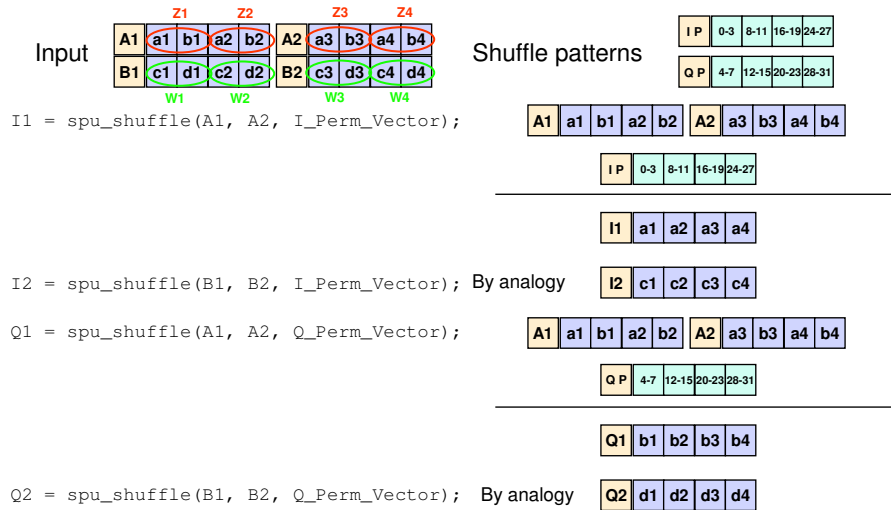
$$(a + ib)(c + id) = (ac - bd) + i(ad + bc)$$

- Or, in code form:

```
/* Given two input arrays with interleaved real and imaginary parts
*/
float input1[2N], input2[2N], output[2N];

for (int i=0;i<N;i+=2) {
    float ac = input1[i]*input2[i];
    float bd = input1[i+1]*input2[i+1];
    output[i] = (ac - bd);
    /*optimized version of (ad+bc) to get rid of a multiply*/
    /* (a+b) * (c+d) -ac - bd = ac + ad + bc + bd -ac -bd = ad + bc */
    output[i+1] = (input1[i]+input1[i+1])*(input2[i]+input2[i+1]) - ac
    - bd;
}
```

Complex Multiplication SPE - Shuffle Vectors



Complex Multiplication

```
A1 = spu_nmsub(Q1, Q2, v_zero);
```

		Q1	b1	b2	b3	b4
*		Q2	d1	d2	d3	d4
-		Z	0	0	0	0
A1	-(b1*d1)	-(b2*d2)	-(b3*d3)	-(b4*d4)		

```
A2 = spu_madd(Q1, I2, v_zero);
```

		Q1	b1	b2	b3	b4
*		I2	c1	c2	c3	c4
+		Z	0	0	0	0
A2	b1*c1	b2*c2	b3*c3	b4*d4		


```
Q1 = spu_madd(I1, Q2, A2);
```

		I1	a1	a2	a3	a4
*		Q2	d1	d2	d3	d4
+		A2	b1*c1	b2*c2	b3*c3	b4*d4
Q1	a1*d1+b1*c1	a2*d2+b2*c2	a3*d3+b3*c3	a4*d4+b4*c4		


```
I1 = spu_madd(I1, I2, A1);
```

		I1	a1	a2	a3	a4
*		I2	c1	c2	c3	c4
+		A1	-(b1*d1)	-(b2*d2)	-(b3*d3)	-(b4*d4)
I1	a1*c1-b1*d1	a2*c2-b2*d2	a3*c3-b3*d3	a4*c4-b4*d4		


Complex Multiplication – Shuffle Back



```
D1 = spu_shuffle(I1, Q1, vcvmrgh);
```



```
D2 = spu_shuffle(I1, Q1, vcvmrgh);
```



Complex Multiplication – SPE - Summary

```

vector float A1, A2, B1, B2, I1, I2, Q1, Q2, D1, D2; /* in-phase (real), quadrature (imag), temp, and
output vectors*/
vector float v_zero = (vector float)(0,0,0,0);

vector unsigned char I_Perm_Vector = (vector unsigned char)(0,1,2,3,8,9,10,11,16,17,18,19,24,25,26,27);
vector unsigned char Q_Perm_Vector = (vector unsigned char)(4,5,6,7,12,13,14,15,20,21,22,23,28,29,30,31);
vector unsigned char vcvmrgh = (vector unsigned char)(0,1,2,3,16,17,18,19,4,5,6,7,20,21,22,23);
vector unsigned char vcvmrgl = (vector unsigned char)(8,9,10,11,24,25,26,27,12,13,14,15,28,29,30,31);

/* input vectors are in interleaved form in A1,A2 and B1,B2 with each input vector representing 2 complex
numbers
and thus this loop would repeat for N/4 iterations */

I1 = spu_shuffle(A1, A2, I_Perm_Vector); /* pulls out 1st and 3rd 4-byte element from vectors A1 and A2 */
I2 = spu_shuffle(B1, B2, I_Perm_Vector); /* pulls out 1st and 3rd 4-byte element from vectors B1 and B2 */
Q1 = spu_shuffle(A1, A2, Q_Perm_Vector); /* pulls out 2nd and 4th 4-byte element from vectors A1 and A2 */
Q2 = spu_shuffle(B1, B2, Q_Perm_Vector); /* pulls out 3rd and 4th 4-byte element from vectors B1 and B2 */
A1 = spu_rmsub(Q1, Q2, v_zero);          /* calculates -(bd - 0) for all four elements */
A2 = spu_madd(Q1, I2, v_zero);           /* calculates (bc + 0) for all four elements */
Q1 = spu_madd(I1, Q2, A2);               /* calculates ad + bc for all four elements */
I1 = spu_madd(I1, I2, A1);               /* calculates ac - bd for all four elements */
D1 = spu_shuffle(I1, Q1, vcvmrgh);       /* spreads the results back into interleaved format */
D2 = spu_shuffle(I1, Q1, vcvmrgl);       /* spreads the results back into interleaved format */

```

SDK Tips and Techniques

- **\$SDK_TOP/make.env contains environment variables to change compiler and compiler settings**
 - SPU_COMPILER
 - tells make to use gcc or xlc for SPU code
 - SPU_TIMING
 - if timing tools RPM is installed will generate a static timing analysis of SPU code to determine pipe stalls and register dependencies
 - SCE_VERSION
 - used to change the toolchain version employed
- **Several internal Systemsim parameters can be accessed and changed via the TCL/TK command line prompt**
 - use the HELP subsystem to access general information about these commands
- **Debugging tools are limited to GDB although there are both PPC and SPU versions**
 - PPC and SPU code must be debugged separately

SDK Tips and Techniques (cont.)


■ Profiling tools

- dynamic profiling
 - `prof_start()`, `prof_stop()`, `prof_clear()`
 - provides branch, pipe utilization, timing, issue, dependency, and other useful performance information during runtime
 - model must be in pipeline mode
- static profiling
 - setting `SPU_TIMING` definition produces a `.timing` file containing a very useful static timing analysis for quickly determining code bottlenecks

Performance Measurement: Profile Checkpoints

NOTE: Execute application on simulator in pipeline mode

```
#include "profile.h"  
prof_cp0(); // clear performance info  
prof_cp30(); // start recording performance info  
< something interesting >  
prof_cp31(); // stop recording performance info
```



(c) Copyright International Business Machines Corporation 2005.
All Rights Reserved. Printed in the United States September 2005.

The following are trademarks of International Business Machines Corporation in the United States, or other countries, or both.
IBM IBM Logo Power Architecture

Other company, product and service names may be trademarks or service marks of others.

All information contained in this document is subject to change without notice. The products described in this document are NOT intended for use in applications such as implantation, life support, or other hazardous uses where malfunction could result in death, bodily injury, or catastrophic property damage. The information contained in this document does not affect or change IBM product specifications or warranties. Nothing in this document shall operate as an express or implied license or indemnity under the intellectual property rights of IBM or third parties. All information contained in this document was obtained in specific environments, and is presented as an illustration. The results obtained in other operating environments may vary.

While the information contained herein is believed to be accurate, such information is preliminary, and should not be relied upon for accuracy or completeness, and no representations or warranties of accuracy or completeness are made.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED ON AN "AS IS" BASIS. In no event will IBM be liable for damages arising directly or indirectly from any use of the information contained in this document.

IBM Microelectronics Division
1580 Route 52, Bldg. 504
Hopewell Junction, NY 12533-6351

The IBM home page is <http://www.ibm.com>
The IBM Microelectronics Division home page is
<http://www.chips.ibm.com>