

Solve Poisson Equations on a Grid

- Goal: explain communication patterns in grid-based problems
 - Domain decomposition
 - Variations on MPI send/receive operations
- Approach:
 - Finite-difference scheme on the Poisson problems
 - Topologies
 - Cartesian topology (decomposition in the x,y,z directions)
 - Use Fortran rather than C

The Poisson Problem

Equations:

$$d^2u / d x^2 + d^2u / d y^2 = f(x,y) \quad (1) \quad \text{in the interior}$$

$$u(x,y) = g(x,y) \quad (2) \quad \text{on the physical boundary}$$

Approximation:


solve equations on a square mesh (grid) *not* everywhere

$$x_i = i/(n+1), \quad i = 0, \dots, n+1,$$

$$y_j = j/(n+1), \quad j = 0, \dots, n+1$$

$$h = 1/(n+1)$$

Five-Point Stencil Approximation



Five-point stencil approximation for 2-D Poisson Problem ($n=7$)

The Poisson Problem (continue)

Finite-difference solution for Eq. (1):

$$u_{i-1,j} + u_{i,j+1} + u_{i,j-1} + u_{i+1,j} - 4u_{i,j} / h^2 = f_{i,j} \quad (3)$$

Use Jacobi iteration to solve for u everywhere on the mesh

Rewrite (3) as

$$u_{i,j} = (u_{i-1,j} + u_{i,j+1} + u_{i,j-1} + u_{i+1,j} - h^2 f_{i,j}) / 4$$

Iterate and update with

$$u^{k+1}_{i,j} = (u^k_{i-1,j} + u^k_{i,j+1} + u^k_{i,j-1} + u^k_{i+1,j} - h^2 f_{i,j}) / 4$$

The Poisson Problem (continue)

The sequential (*single* process) code:

```
integer i, j, nx, ny
```

```
double precision u(0:nx+1,0:ny+1), unew(0:nx+1,0:ny+1)
```


```
do 10 j=1, ny
```

```
  do 10 i=1, nx
```

```
    unew(i,j) = 0.25 * (u(i-1,j)+u(i,j+1)+u(i,j-1)+u(i+1,j)) - h * h * f(i,j)
```

```
  10 continue
```

1D Domain Decomposition



1D domain decomposition along the y direction

The Poisson Problem (continue)

The code on *each* process with *1D domain decomposition* along *y/j*:


```
integer i, j, nx, ny
C s:e indicates the value of j that this process is responsible for.
double precision u(0:nx+1, s:e), unew(0:nx+1, s:e)

do 10 j=s,e
  do 10 i=1, nx

    unew(i,j) = 0.25 * (u(i-1,j)+u(i,j+1)+u(i,j-1)+u(i+1,j)) - h * h * f(i,j)

10 continue
```

1D Domain Decomposition



The computational domain, with 1-point ghost zone, for one of the processes

The Poisson Problem (continue)

The code on each process with 1D domain decomposition along y/j
with *1-point ghost zone*:


```
integer i, j, nx, ny
C need grid points from top and bottom processes.
  double precision u(0:nx+1, s-1:e+1), unew(0:nx+1, s-1:e+1)

C the do-loop bound is still from s to e along the y direction
  do 10 j=s,e
    do 10 i=1, nx
      unew(i,j) = 0.25 * (u(i-1,j)+u(i,j+1)+u(i,j-1)+u(i+1,j)) - h * h * f(i,j)
    10 continue
```

Topologies

- Topologies is the description of how the processes in a parallel computer are connected to one another (or more precisely, of the **interconnection network**).
- How a particular application **map** to the physical topology of the parallel computer depends on the **vendor**
- One MPI application topology is the **Cartesian** topology which is simply a decomposition in the x,y,z directions

2D Cartesian Decomposition



Note: The coordinate is (column, row) which would be returned by `MPI_Get_coords`. The rank is returned by `MPI_Cart_Rank`.

2D Cartesian Decomposition (code)


```
dims(1)      = 4
dims(2)      = 3
periods(1)   = .false.
periods(2)   = .false.
reorder      = .true.
ndim         = 2
call MPI_CART_CREATE( MPI_COMM_WORLD, ndim,
                     dims, periods, reorder, comm2d, ierr)
```

Note: *reorder* = *.true.* allows MPI to map the processes to the elements of the decomposition optimally



MPI_Cart_Shift

*Int MPI_Cart_shift(MPI_Comm comm, int direction, int displ, int *src, int*dest)*


MPI_Cart_Shift() is used to find the neighbors. For example, shift by one in the 1st dimension. (1,1)'s destination and source are (2,1) and (0,1), respectively



Direction=0
Displ=1



1D Domain Decomposition (Transfer Data)



Main Program

```
c*****
c oned.f - a solution to the Poisson problem using Jacobi
c iteration on a 1-d decomposition
c
C The size of the domain is read by processor 0 and broadcast to
c all other processors. The Jacobi iteration is run until the
c change in successive elements is small or a maximum number of
c iterations is reached. The difference is printed out at each step.
c*****
program main

include "mpif.h"
integer maxn parameter (maxn = 128)
double precision a(maxn,maxn), b(maxn,maxn), f(maxn,maxn)
integer nx, ny
```

```

integer myid, numprocs, ierr
integer comm1d, nbrbottom, nbrtop, s, e, it, maxit
double precision diff, diffnorm, dwork
double precision t1, t2
double precision MPI_WTIME
external MPI_WTIME
external diff

call MPI_INIT( ierr )
call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )
call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )

if (myid .eq. 0) then
c
c Get the size of the problem
c
  print *, 'Enter nx'
  read *, nx
endif

```



```

call MPI_BCAST(nx,1,MPI_INTEGER,0,MPI_COMM_WORLD,ierr)
ny = nx
C
c Get a new communicator, comm1d, for a 1D decomposition of the domain
C
call MPI_CART_CREATE( MPI_COMM_WORLD, 1, numprocs,
$                      .false., .true., comm1d, ierr )
C
c Get my position, myid, in this communicator, comm1d, and my neighbors
c at the top, nbrtop and the bottom, nbrbottom
C
call MPI_COMM_RANK( comm1d, myid, ierr )
call MPI_Cart_shift( comm1d, 0, 1, nbrbottom, nbrtop, ierr )
C
c Compute the actual decomposition: s and e
call MPE_DECOMP1D( ny, numprocs, myid, s, e )
C

```

c Initialize the right-hand-side (f) and the initial solution guess (a)

c

call onedinit(a, b, f, nx, s, e)

c

c Actually do the computation. Note the use of a collective operation to

c check for convergence, and a do-loop to bound the number of iterations.

c

call MPI_BARRIER(MPI_COMM_WORLD, ierr)

t1 = MPI_WTIME()

do 10 it=1, maxit

C get ghost points

call exchng1(a, nx, s, e, comm1d, nbrbottom, nbrtop)

c perform one jacobi “sweep”

call sweep1d(a, f, nx, s, e, b)

call exchng1(b, nx, s, e, comm1d, nbrbottom, nbrtop)

call sweep1d(b, f, nx, s, e, a)

dwork = diff(a, b, nx, s, e)

call MPI_Allreduce(dwork, diffnorm, 1, MPI_DOUBLE_PRECISION,
\$ MPI_SUM, comm1d, ierr)

```

        if (diffnorm .lt. 1.0e-5) goto 20
c
        if (myid .eq. 0) print *, 2*it, ' Difference is ', diffnorm
10    continue


        if (myid .eq. 0)
        print *, 'Failed to converge'

20    continue
    t2 = MPI_WTIME()

        if (myid .eq. 0) then
        print *, 'Converged after ', 2*it, ' Iterations in ', t2 - t1,
        $ ' secs ' endif
c
        call MPI_FINALIZE(ierr)
        end

```

1D Domain Decomposition (Transfer Data)



Exchange Data for Ghost Points (I)

Using blocking sends and receives

```
subroutine exchn1( a, nx, s, e, comm1d, nbrbottom, nbrtop )  
include 'mpif.h'  
integer nx, s, e  
double precision a(0:nx+1,s-1:e+1)  
integer comm1d, nbrbottom, nbrtop  
integer status(MPI_STATUS_SIZE), ierr
```

C

```
call MPI_SEND( a(1,e), nx, MPI_DOUBLE_PRECISION, nbrtop, 0, comm1d, ierr)  
call MPI_RECV(a(1,s-1), nx, MPI_DOUBLE_PRECISION, nbrbottom, 0,  
& comm1d, status, ierr )  
  
call MPI_SEND(a(1,s), nx, MPI_DOUBLE_PRECISION, nbrbottom, 1, comm1d,ierr)  
call MPI_RECV(a(1,e+1), nx, MPI_DOUBLE_PRECISION, nbrtop, 1,  
& comm1d, status, ierr )  
return  
end
```

Exchange Data for Ghost Points (II)

Using **paired** blocking sends and receives

```
...
integer comm1d, nbrbottom, nbrtop, rank, coord
...
call MPI_COMM_RANK(comm1d, rank, ierr)
call MPI_CART_COORDS(comm1d, rank, 1, coord, ierr)
if (mod(coord, 2) .eq. 0) then
  call MPI_SEND( a(1,e), nx, MPI_DOUBLE_PRECISION, nbrtop, 0, comm1d, ierr)
  call MPI_RECV(a(1,s-1), nx, MPI_DOUBLE_PRECISION, nbrbottom, 0,
    &
    comm1d, status, ierr )
  call MPI_SEND(a(1,s), nx, MPI_DOUBLE_PRECISION, nbrbottom, 1, comm1d,ierr)
  call MPI_RECV(a(1,e+1), nx, MPI_DOUBLE_PRECISION, nbrtop, 1,
    &
    comm1d, status, ierr )
else
  call MPI_RECV(a(1,s-1), nx, MPI_DOUBLE_PRECISION, nbrbottom, 0,
    &
    comm1d, status, ierr )
  call MPI_SEND( a(1,e), nx, MPI_DOUBLE_PRECISION, nbrtop, 0, comm1d, ierr)
  call MPI_RECV(a(1,e+1), nx, MPI_DOUBLE_PRECISION, nbrtop, 1,
    &
    comm1d, status, ierr )
  call MPI_SEND(a(1,s), nx, MPI_DOUBLE_PRECISION, nbrbottom, 1, comm1d,ierr)
...

```

Exchange Data for Ghost Points (III)

Using send-recv

```
subroutine exchng1( a, nx, s, e, comm1d, nbrbottom, nbrtop )  
include 'mpif.h'  
integer nx, s, e  
double precision a(0:nx+1,s-1:e+1)  
integer comm1d, nbrbottom, nbrtop  
integer status(MPI_STATUS_SIZE), ierr
```

C

```
call MPI_SENDRECV(  
&          a(1,e), nx, MPI_DOUBLE_PRECISION, nbrtop, 0,  
&          a(1,s-1), nx, MPI_DOUBLE_PRECISION, nbrbottom, 0,  
&          comm1d, status, ierr )  
  
call MPI_SENDRECV( & a(1,s), nx, MPI_DOUBLE_PRECISION, nbrbottom, 1,  
&          a(1,e+1), nx, MPI_DOUBLE_PRECISION, nbrtop, 1,  
&          comm1d, status, ierr )  
return  
end
```

Exchange Data for Ghost Points (IV)

Using nonblocking operations

```
subroutine exchn1( a, nx, s, e, comm1d, nbrbottom, nbrtop )  
include 'mpif.h'  
integer nx, s, e  
double precision a(0:nx+1,s-1:e+1)  
integer comm1d, nbrbottom, nbrtop  
integer status_array(MPI_STATUS_SIZE,4), ierr, req(4)
```

C

```
call MPI_Irecv(a(1,s-1), nx, MPI_DOUBLE_PRECISION, nbrbottom, 0,  
&               comm1d, req(1), ierr )  
call MPI_Irecv(a(1,e+1), nx, MPI_DOUBLE_PRECISION, nbrtop, 1,  
&               comm1d, req(2), ierr )  
  
call MPI_Isend( a(1,e), nx, MPI_DOUBLE_PRECISION, nbrtop, 0,  
&               comm1d, req(3), ierr )  
call MPI_Isend(a(1,s), nx, MPI_DOUBLE_PRECISION, nbrbottom, 1,  
&               comm1d, req(4), ierr )  
return  
end
```


Compute Decomposition

```
c This file contains a routine for producing a decomposition of a 1-d array
c when given a number of processors. It may be used in "direct" product
c decomposition. The values returned assume a "global" domain in [1:n]
c n, numprocs, myid are input while s and e are output
  subroutine MPE_DECOMP1D( n, numprocs, myid, s, e )
    integer n, numprocs, myid, s, e, nlocal, deficit
c
    nlocal = n / numprocs
    s = myid * nlocal + 1
    deficit = mod(n,numprocs)
    s = s + min(myid,deficit)
    if (myid .lt. deficit) then
      nlocal = nlocal + 1
    endif e = s + nlocal - 1
    if (e .gt. n .or. myid .eq. numprocs-1) e = n
    return
  end
```

Calculate the Difference of Two Successive Solutions

```
double precision function diff( a, b, nx, s, e )  
integer nx, s, e  
double precision a(0:nx+1, s-1:e+1), b(0:nx+1, s-1:e+1)  
c  
double precision sum integer i, j  
c  
sum = 0.0d0  
do 10 j=s,e  
do 10 i=1,nx  
sum = sum + (a(i,j) - b(i,j)) ** 2  
10 continue  
c  
diff = sum  
return  
end
```

Initialization Routine

```
subroutine onedinit( a, b, f, nx, s, e )
```

```
integer nx, s, e, i,j
```

```
double precision a(0:nx+1, s-1:e+1), b(0:nx+1, s-1:e+1),
```

```
&                f(0:nx+1, s-1:e+1)
```

```
c
```

```
do 10 j=s-1,e+1
```

```
do 10 i=0,nx+1
```

```
a(i,j) = 0.0d0
```

```
b(i,j) = 0.0d0
```

```
f(i,j) = 0.0d0
```

```
10 continue
```

c Handle boundary conditions

do 20 j=s,e

a(0,j) = 1.0d0

b(0,j) = 1.0d0

a(nx+1,j) = 0.0d0

b(nx+1,j) = 0.0d0

20 continue

if (s .eq. 1) then

do 30 i=1,nx

a(i,0) = 1.0d0

b(i,0) = 1.0d0

30 continue

endif

return

end

Jacobi Sweep

c Perform a Jacobi sweep for a 1-d decomposition.

c Sweep from a into b

c

subroutine sweep1d(a, f, nx, s, e, b)

integer nx, s, e, i,j

double precision a(0:nx+1,s-1:e+1), f(0:nx+1,s-1:e+1),
+ b(0:nx+1,s-1:e+1), h

```

h = 1.0d0 / dble(nx+1)
do 10 j=s, e
  do 10 i=1, nx
    b(i,j) = 0.25 * (a(i-1,j)+a(i,j+1)+a(i,j-1)+a(i+1,j)) -
$          h * h * f(i,j)
10 continue
return
end

```

#3 Homework for MPI

- In the header, list your name, email address, general description of the code. Inside code, add sufficient comments on each nontrivial section

3 Homework for MPI

– Solve a 2D wave equation with a finite-difference scheme

- Wave equation: $-\phi_{,tt}/c^2 + \phi_{,xx} + \phi_{,yy} + \phi_{,zz} = 0$
- Explicit *finite differencing* (centered, second order)

$$\begin{aligned}\phi^{n+1}_{i,j} = & 2\phi^n_{i,j} - \phi^{n-1}_{i,j} \\ & + \Delta t^2/\Delta x^2 (\phi^n_{i+1,j} - 2\phi^n_{i,j} + \phi^n_{i-1,j}) \\ & + \Delta t^2/\Delta y^2 (\phi^n_{i,j+1} - 2\phi^n_{i,j} + \phi^n_{i,j-1})\end{aligned}$$

- Initial condition: $\phi(i,j) = \sin(i \cdot dx) \cdot \cos(j \cdot dy)$, where $dx = 2\pi/\text{GridSizeX}$, $dy = 2\pi/\text{GridSizeY}$.
- Free boundary condition along the x direction and periodic along the y direction
- Parameters: $\text{GridSizeX}=512$, $\text{GridSizeY}=1024$, $c=1$; use 2D domain decomposition with 2 and 4 processes in x and y directions, respectively; run 10 time steps; i is from 0 to 511 and j is from 0 to 1023.
- Compare the parallel code with the serial code in performance and accuracy. You are encouraged to vary the number of grid points along x and y to analyze the performance scalability.